

CPSC 131

Data Structures Concepts

Unordered Containers / Hash Tables

Dr. Anand Panangadan

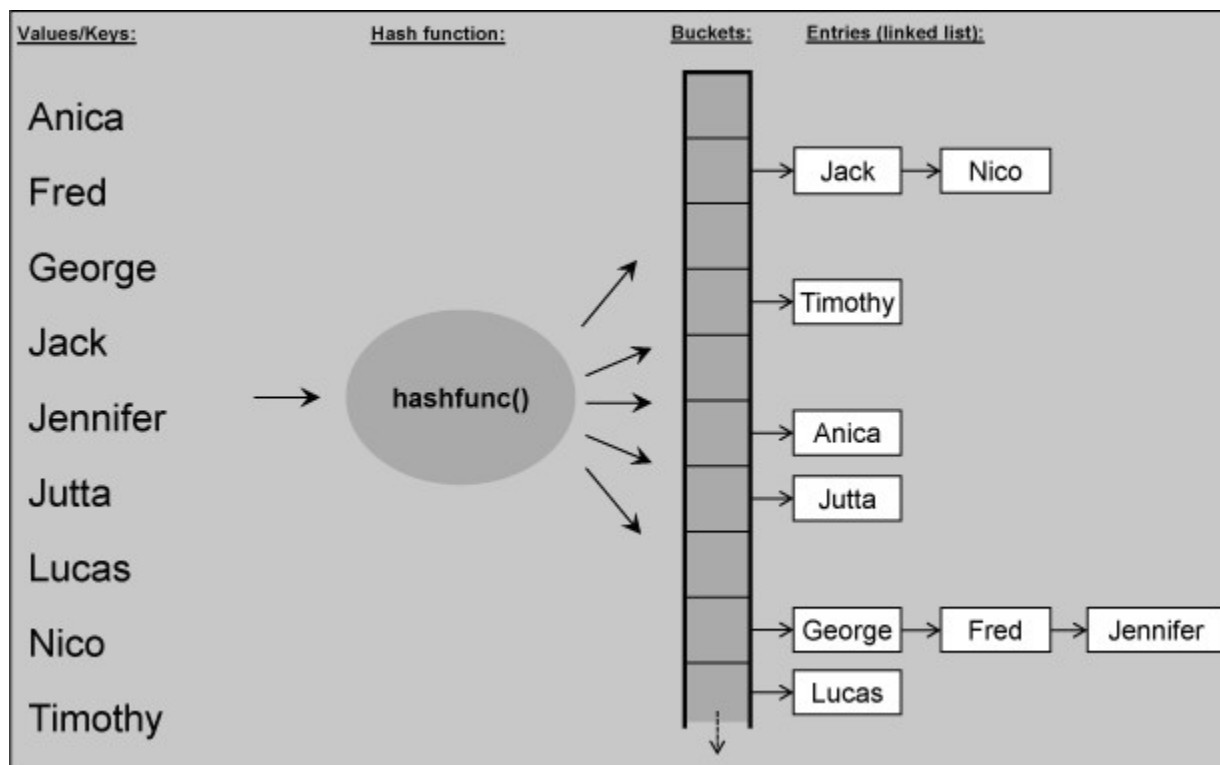
apanangadan@fullerton.edu

Dr. Thomas Bettens

TLBettens@fullerton.edu

Unordered Containers / Hash Tables

Josuttis, *The C++ Standard Library*



Elements have no defined order

Finding an element is faster than associative containers

Unordered Containers

*Josuttis, The C++
Standard Library*

- In unordered containers, elements have no defined order
 - If you insert three elements, they might have any order when you iterate over all the elements in the container.
 - If you insert a fourth element, the order of the elements previously inserted might change.
 - The only important fact is that a specific element is somewhere in the container.
 - Even when you have two containers with equal elements inside, the order might be different.
 - Think of it as like a bag.

Unordered Containers

*Josuttis, The C++
Standard Library*

- Unordered containers are typically implemented as a hash table.
 - Internally, the container is an array of linked lists.
- Using a hash function, the position of an element in the array gets processed.
 - The goal is that each element has its own position so that you have fast access to each element, provided that the hash function is fast.
 - Multiple elements might have the same position because such a fast perfect hash function is not always possible or might require that the array consumes a huge amount of memory.
 - For this reason, the elements in the array are linked lists so that you can store more than one element at each array position.

Unordered Containers


*Josuttis, The C++
Standard Library*

Major advantage of unordered containers


- Finding an element with a specific value is even faster than for associative containers.
 - The use of unordered containers provides amortized constant complexity, provided that you have a good hash function.
 - However, providing a good hash function is not easy

A constant time data structure?

- How long does it take to find an entry in a data structure?



- Linked lists: $O(n)$
- Balanced trees (AVL): $O(\log n)$



- Can we get to $O(1)$?
- An array can get values in $O(1)$ *if*
 - Keys are the same as array index




A constant time data structure?

- An array can get values in $O(1)$ *if*
 - Keys are the same as array index
- Disadvantages:
 - Requires keys be unique integers in the range $0, 1, \dots, N-1$
 - wastes a lot of space if the number of entries are much smaller than N
- The hash table is an attempt to reach $O(1)$ by:
 - converting keys into codes, which may not be unique
 - compressing codes into indexes within a reduced storage space

Operation vs Time Summary


Windows 32 bit x86 image (executed in Safe Mode)



Main idea

Convert any data type into
an array index

Hashing function




Hashtable

- A **hash table** for a given key type consists of
 - Hash function h
 - Array (called **table**) of size N
- When implementing a map with a hash table, the goal is to store item (k, o) at index $i = h(k)$

(Key, value) pairs

- Reminder:
 - Keys are associated with values
 - For simplicity, only showing keys in figures



Compression function

Division method


- Let N be the size of the array
- To get array index of hash code k , do
- **$\text{Index} = k \% N$**
 - Take remainder after dividing k by N
- Simple and commonly used

Example:

- Let $N=7$
- Key 20 goes into array index
 $20 \% 7 = 6$
 - $\text{array}[6] = 20$

Compression function

Insert keys 20, 25, 30 into a table of size 7




Compression function

Searching is similar to insert
search(25):


1. Calculate $\text{index}(25) = 25\%7 = 4$
2. Look in `array[4]`

Cost of insert/search = $O(1)$



Compression function

Inserted keys 20, 25, 30 into a table of size 7
Insert key 16

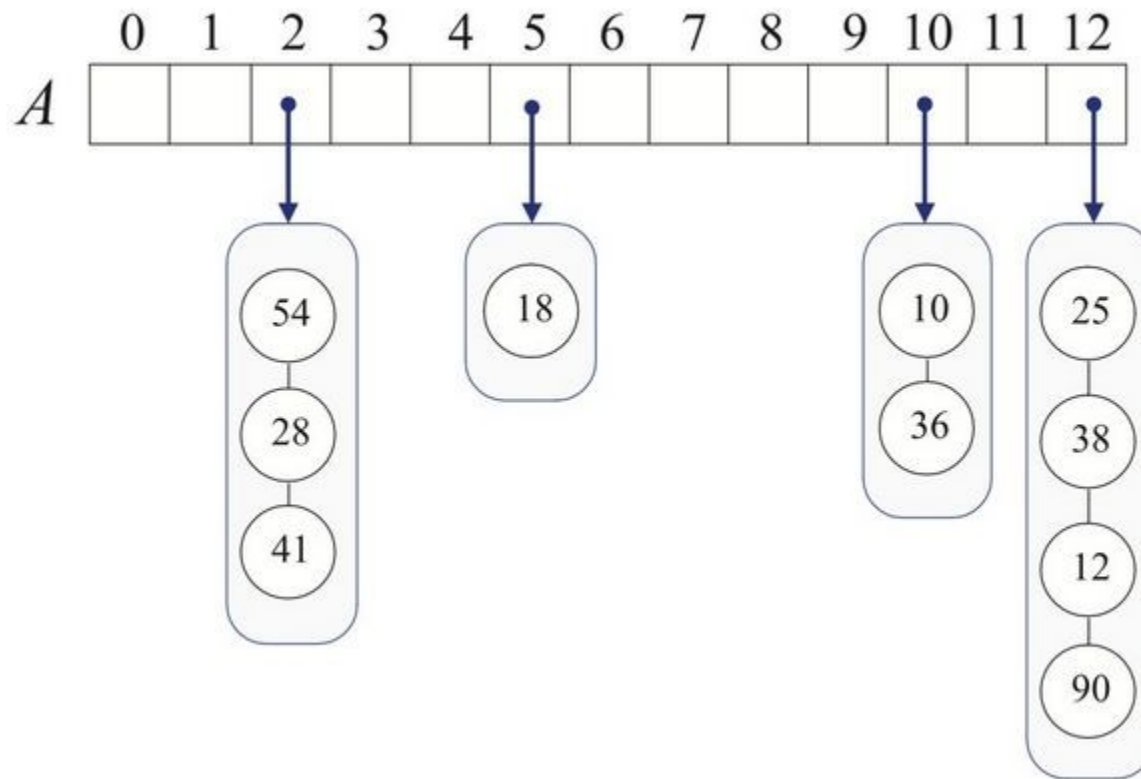


Collisions

- When two *different* keys get assigned to the *same* table index
 - $30 \% 7 = 16 \% 7 = 2$
- Solutions to deal with collisions
 - Chaining
 - Probing
 - Linear probing
 - Quadratic probing
 - Prevent collisions completely – *Direct Hashing*

Chaining

Have each bucket hold a **list** (or a **vector**) of keys



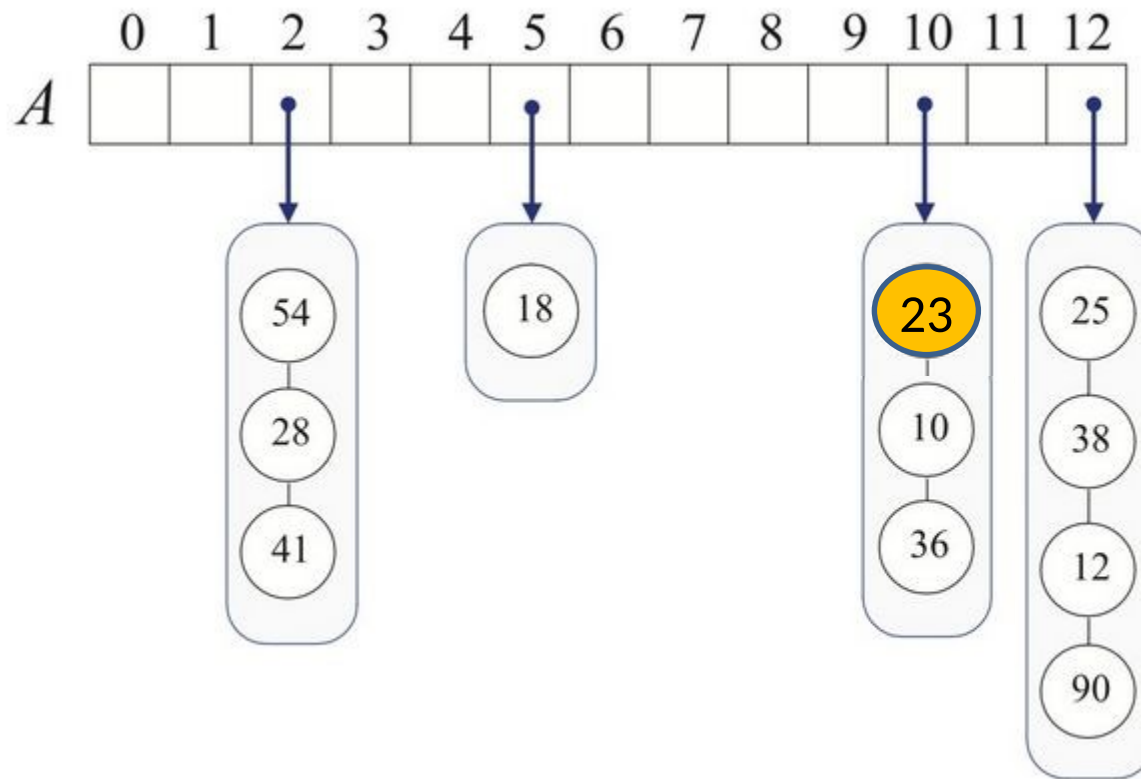
Chaining

Have each bucket hold a **list** (or a **vector**) of keys

- find(): calculate hash, search bucket's list for key
- insert(): calculate hash, insert key into bucket's list.
- remove(): calculate hash, remove key from bucket's list

Chaining

Insert(**23**)? *Note that N=13*



Chaining

Have each bucket hold a **list** (or a **vector**) of keys


- `find()`: calculate hash, search bucket's list for key
 - $O(n)$ **worst-case**
- `insert()`: calculate hash, insert key into bucket's list.
 - $O(1)$ **worst-case if duplicates allowed**
 - $O(n)$ **worst-case if duplicates not allowed**
- `remove()`: calculate hash, remove key from bucket's list
 - $O(n)$ **worst-case**

Linear probing

- Only have a single array
- If bucket is occupied, search forward for a free bucket
- Search is circular
 - when end of table is reached, wrap around to beginning
- Search fails if starting point is reached

Resolving collisions with linear probing

A hash table of size 13, compression function: $\text{key} \% 13$



Entries: 18 41 22 44 59 32 31 73
key mod 13: 5 2 9 5 7 6 5 8

0	1	2	3	4	5	6	7	8	9	10	11	12
		41			18	44	59		22			

Values of “empty” cells

- Two kinds of empty cells
 - Empty since start (E1)
 - Empty after removal of an element (E2)
- Initially, all cells have value E1

Search with Linear Probing

- Consider a hash table A that uses linear probing
- $\text{find}(k)$
 - We start at cell $h(k)$
 - We probe consecutive locations until one of the following occurs
 - An item with key k is found, or
 - An empty cell (E1) is found, or
 - N cells have been unsuccessfully probed

Algorithm $\text{find}(k)$

$i \leftarrow h(k)$

$n \leftarrow 0$

repeat

$c \leftarrow A[i]$

if $c == \text{E1}$

return

not found

else if $c.\text{key}()$

$== k$

return

$c.\text{value}()$

else

$i \leftarrow (i +$

$1) \bmod N$

$n \leftarrow n + 1$

until

$n == N$

return *not found*

Insert with Linear Probing

- $\text{insert}(k, \text{value})$
 - We start at cell $h(k)$
 - We probe consecutive locations until an empty cell (E1 or E2) is found, or
 - N cells have been unsuccessfully probed

Algorithm $\text{insert}(k, \text{value})$

```
 $i \leftarrow h(k)$   
 $n \leftarrow 0$   
repeat  
     $c \leftarrow A[i]$   
    if  $c == \text{E1}$  or  
    E2  $A[i].\text{key} \leftarrow k$   
     $A[i].\text{value} \leftarrow \text{value}$   
    break  
    else  $i \leftarrow (i + 1) \bmod N$   
     $n \leftarrow n + 1$   
until  $n == N$ 
```

Remove with Linear Probing

- **remove(k)**
 - We start at cell $h(k)$
 - We probe consecutive locations until one of the following occurs
 - An item with key k is found, or
 - An empty cell (E1) is found, or
 - N cells have been unsuccessfully probed

Algorithm *remove(k)*

$i \leftarrow h(k)$

$n \leftarrow 0$

repeat

$c \leftarrow A[i]$

if $c == E1$

return

not found

else if $c.key ()$

$== k$

$A[i] \leftarrow$

E2

return

else

$i \leftarrow (i +$

$1) \bmod N$

$n \leftarrow n + 1$

until


$n == N$

Performance of Linear Probing

- Colliding items lump together, causing future collisions to cause a longer sequence of probes
- In the **worst case**, searches, insertions and removals on a hash table take **$O(n)$** time
 - The worst case occurs when all the keys inserted into the map collide
- **Load factor** of a hash table $\alpha = n/N$
 - How full is the hash table?
- The load factor affects the performance of a hash table
- Assuming that the hash values are like random numbers, the **expected** number of probes for an insertion with linear probing is
$$1 / (1 - \alpha)$$
- Recommendation: keep $\alpha < 0.5$
 - At least half the table must be empty
- Then, **expected cost** $< 1/(1-0.5) = O(1)$

Performance of Hashing

- The **expected** running time of all the operations in a hash table is **$O(1)$**
- In practice, **hashing is very fast provided the load factor is not close to 100%**
- Recommendations:
 - keep $\alpha < 0.5$ for linear probing
 - keep $\alpha < 0.9$ for separate chaining
- Applications of hash tables:
 - small databases
 - compilers
 - browser caches



Quadratic probing


- Same approach as linear probing but probe sequence is different
- Linear probing sequence ():
 - $\text{index} = H, H+1, H+2, H+3, \dots$
- Quadratic probing sequence:
 - c_1 and c_2 are constants that are given
 - For instance: $c_1=1, c_2=1$

Quadratic probing

- Example:
 - Let table size
- Probe sequence for key=45:
 - First probe:
 - Then probe:
 - Then probe:


Direct Hashing

- Can prevent collisions **completely**, if:
 - Table is large enough that every key gets its own index
 - Array index = key
 - A true $O(1)$ data structure




- Limitations of direct hashing
 1. All keys must be non-negative integers
 2. The hash table's size equals the largest key value plus 1, which may be very large

Hash Codes



- How do we deal with non-integer keys?
 - Char
 - Strings
- The goal of a hash function is to “disperse” the keys in an apparently random way

Hash Codes




- ASCII code
 - A number for every character
 - “A” = 65; “B” = 66, “C” = 67, ..., “Z”=90
 - “a” = 97; “b” = 98, “c” = 99, ..., “z”=122

ASCII table


Dec	Char	Dec	Char	Dec	Char
32	[space]	64	@	96	`
33	!	65	A	97	a
34	"	66	B	98	b
35	#	67	C	99	c
36	\$	68	D	100	d
37	%	69	E	101	e
38	&	70	F	102	f
39	'	71	G	103	g
40	(72	H	104	h
41)	73	I	105	i
42	*	74	J	106	j
43	+	75	K	107	k
44	,	76	L	108	l
45	-	77	M	109	m
46	.	78	N	110	n
47	/	79	O	111	o
48	0	80	P	112	p
49	1	81	Q	113	q
50	2	82	R	114	r
51	3	83	S	115	s
52	4	84	T	116	t
53	5	85	U	117	u
54	6	86	V	118	v
55	7	87	W	119	w
56	8	88	X	120	x
57	9	89	Y	121	y
58	:	90	Z	122	z
59	;	91	[123	{
60	<	92	\	124	
61	=	93]	125	}
62	>	94	^	126	~
63	?	95	_	127	

Hash Codes




- Component sum
 - If the key is longer than an integer
 - Partition the key into components of fixed length (e.g., 16 or 32 bits) and sum the components (ignoring overflows)
 - Can we do this for strings?
 - “A” = 65
 - “AB” = 65 + 66 = 131
 - “ABE” = 65 + 66 + 69 = 200

Hash Codes



- Component sum for strings
 - “A” = 65
 - “AB” = 65 + 66 = 131
 - “ABE” = 65 + 66 + 69 = 200
- What is the issue?
 - Order does **not** matter
 - “RAMON” and “NORMA” have the same hash code

Hash codes



Polynomial accumulation

- Again, split key into components (x_0, x_1, \dots)
- But instead of just adding,
- Treat components as the coefficients of a polynomial:

$$x_0 a^{k-1} + x_1 a^{k-2} + \dots + x_{k-2} a + x_{k-1}$$


“a” is some number

- Easier to write code when rewritten like this:

$$x_{k-1} + a(x_{k-2} + a(x_{k-3} + \dots a(x_2 + a(x_1 + a x_0)) \dots))$$

N	O	R	M	A
x_0	x_1	x_2	x_3	x_4


Hash codes



C++ implementation of Polynomial accumulation

```
size_t polynomial_hash (string word) {  
    const int a = 33;  
    size_t hash = 0;  
    for (int i = 0; i < word.size(); i++)  
        hash = hash*a + word[i];  
    return hash;  
}
```


Hash codes



Polynomial accumulation

- Especially suitable for strings
- Choice of $a = 33$ gives at most 6 collisions on a set of 50,000 English words

Hash Codes



- Integer cast
 - We reinterpret the bits of the key as an integer
 - Suitable for keys of length less than or equal to the number of bits of the integer type
 - byte, short, int, float in C++

hash table in C++

- Usage similar to a map

```
#include <unordered_map>
```

```
std::unordered_map<std::string, double> gpaRecord;
```

```
gpaRecord["Allen"] = 3.42;    // new element inserted  
gpaRecord["Beth"] = 3.5;     // new element inserted  
cout << gpaRecord["Allen"];  // existing element read
```

- “**unordered_map** containers are **faster** than **map** containers to access individual elements by their key, although they are generally **less efficient for range iteration** through a subset of their elements.”

hash table in C++

- `std::unordered_map`
- Using the operator[key] automatically inserts (key, default value) if key is not found!
 - Same as in `std::map`
 - Check for key using `find()` and end iterator

```
std::unordered_map<std::string, double> gpaRecord;  
  
if( gpaRecord.find("Allen") != gpaRecord.end() )  
{  
    cout << gpaRecord["Allen"];  
}
```